


Dimensionality
reduction

BIPN 162



By the end of
today's lecture you
will be able to:

- Motivate the use of dimensionality reduction in neuroscience
- Identify several popular methods and their limitations
- Define what eigenvalues and eigenvectors are and determine them using Python
- Describe the steps of PCA and implement it in Python
 - Define and generate a **covariance matrix**



**Dimensionality
reduction** assumes
that the messy &
complex world comes
from the interplay of a
few “latent” factors

“Describing a complex signal,
such as a visual scene or a pattern
of neural activity in terms of just a
few summarizing features”

- Pang et al. 2016

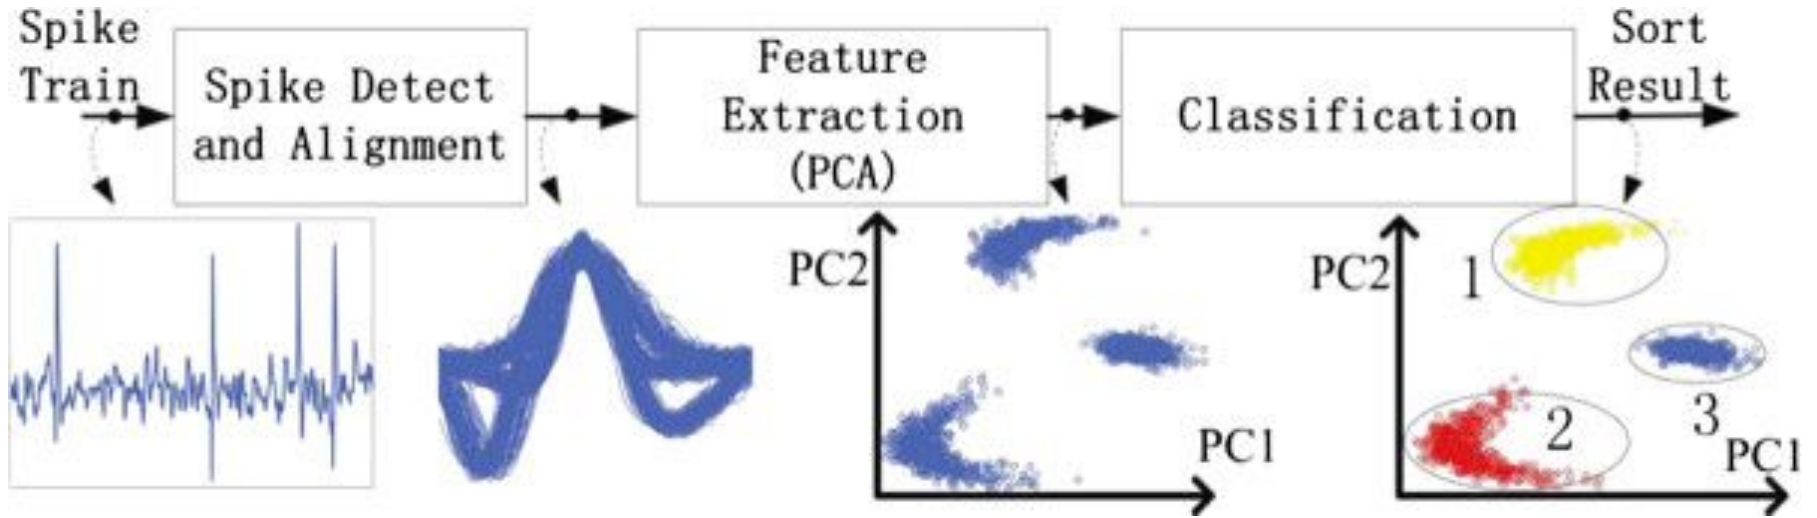
Implementations of dimensionality reduction *beyond* neuroscience

- Personality (five-factor OCEAN model; Costa & McCrae, 1992)
- Facial recognition
- Handwriting recognition

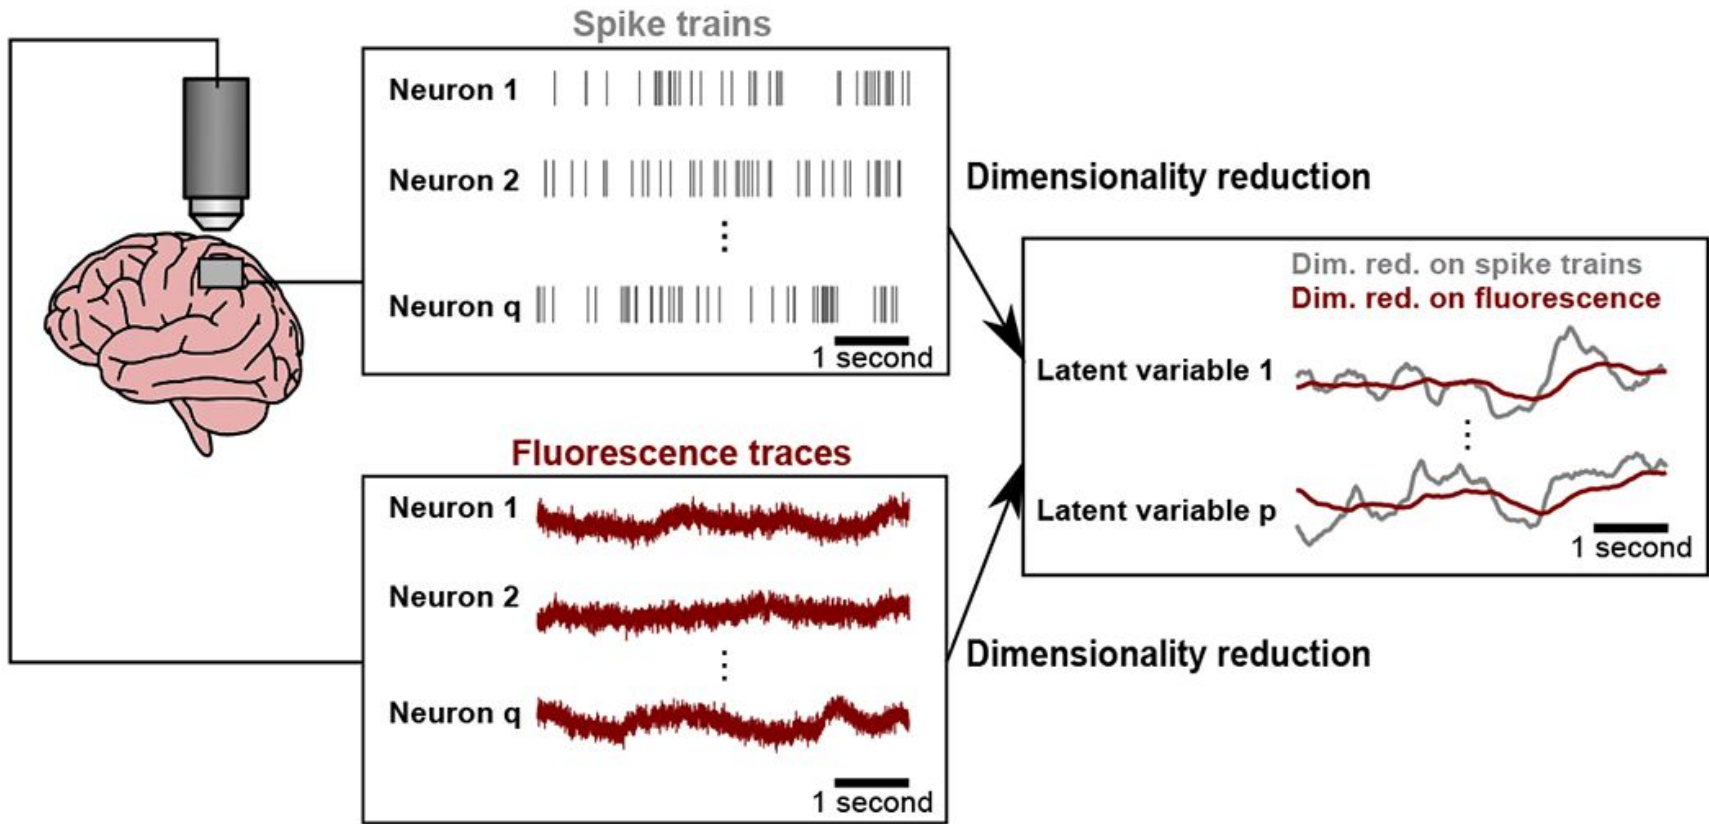


Eigenfaces; try it yourself

Image source




We've seen dimensionality reduction before...



The use of dimensionality reduction on spike trains or fluorescence

<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10358781/>

Large-scale datasets + Data-driven modeling




We've seen dimensionality reduction before... (slide: Dr. Khosla)

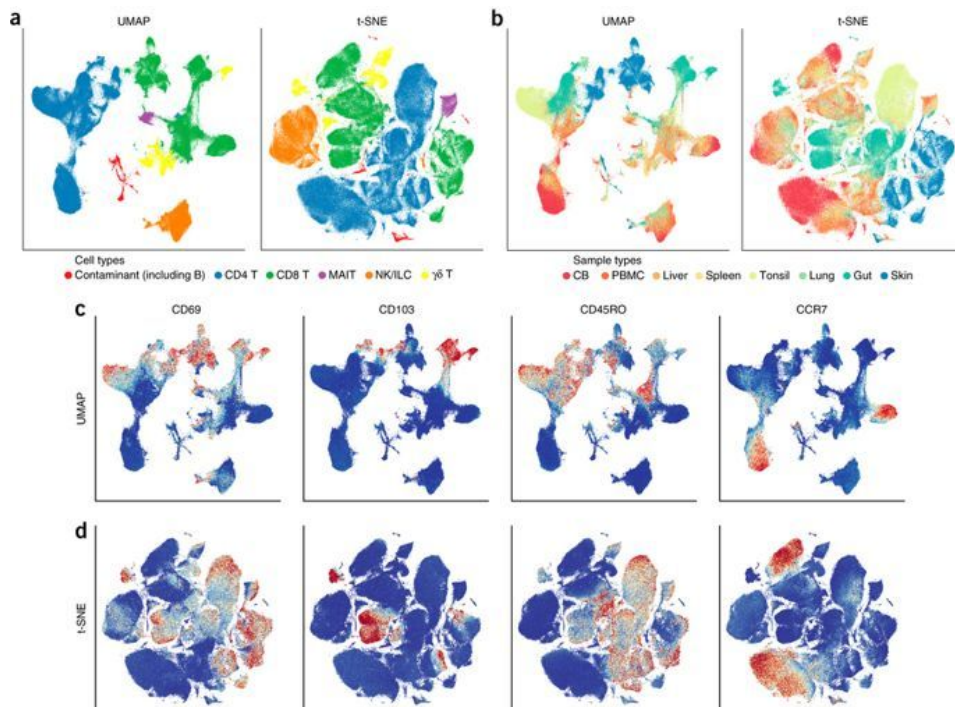
Matrix factorization

Often **non-negative matrix factorization (NMF)**, which enforces positive values.

A simple and general framework for extracting correlations and low-dimensional structure from matrix-coded datasets



Implementations of dimensionality reduction



Single-cell techniques ([paper](#))

t-SNE:

<https://lvdmaaten.github.io/tsne/>

Why do we need dimensionality reduction in neuroscience?

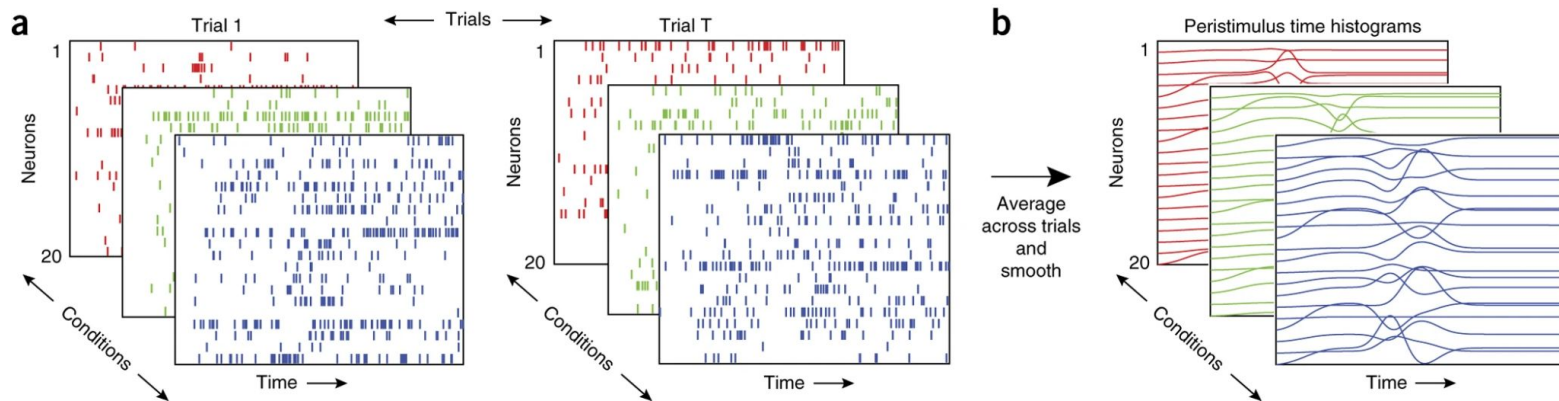
- Neural systems and behaviors are complex
 - We need many numbers to characterize systems like networks of neurons or even a single neuron — in other words, our data is **multivariate**
- Dimensionality reduction can help reduce the complexity of systems/data


“[W]e ought to seek representations, both of the stimulus and of brain activity, that are **concise**, **complete**, and **informative** about the workings of the nervous system, and yet which are not biased by an experimenter’s arbitrary choice. **Considering this task from the perspective of dimensionality reduction provides an entry point into principled mathematical techniques that let us discover these representations directly from experimental data**, a key step to developing rich yet comprehensible models for brain function.”

Pang et al. (2016)

Three primary uses of dimensionality reduction for large-scale recordings in neuroscience:

1. Retaining population structure for single-trial analyses
2. Extracting hidden “latent” features from data to test hypotheses about neural population structure
3. Exploratory data analysis








Population activity r_1 , r_2 and r_3 can be reconstructed by taking a weighted combination of the latent variables, where the weights are specified by the matrix shown.

Population activity (black points) lies in a plane (shaded gray). Each point represents the population activity at a particular time and can be equivalently referred to using its high-dimensional coordinates $[r_1, r_2, r_3]$ or low-dimensional coordinates $[s_1, s_2]$. The points trace out a trajectory over time (black curve).

The latent variables s_1 and s_2 can be obtained by taking a weighted combination of the population activity, where the weights are specified by the matrix shown.



PCA and K-Means identified six different clusters of neurons, which are also spatially-organized (right). These clusters are similar during both spontaneous and odor-evoked activity.

Report

Spontaneous Activity Governs Olfactory Representations in Spatially Organized Habenular Microcircuits

Suresh Kumar Jetti,^{1,2,4} Nuria Vendrell,¹ and Emre Yaksi^{1,2,3,*}

¹NERF, Kapeldreef 75, 3001 Leuven, Belgium

²KU Leuven, Kapeldreef 75, 3001 Leuven, Belgium

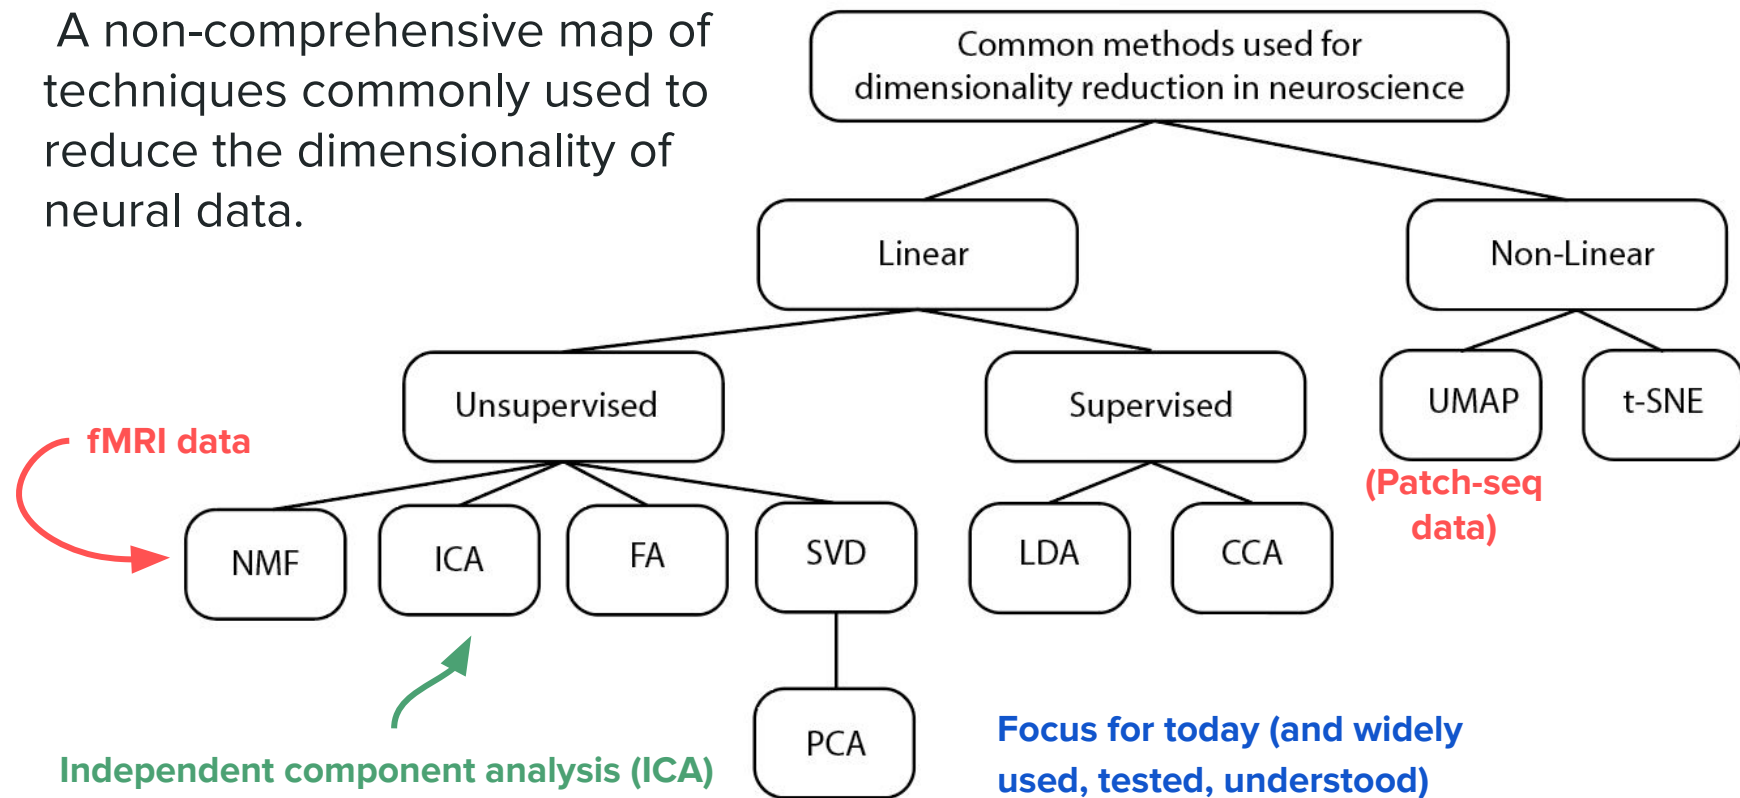
³VIB, Kapeldreef 75, 3001 Leuven, Belgium

Moreover, we show that the spontaneous dHb activity is not random but structured into functionally and spatially organized clusters of neurons, which reflects the favored states of the dHb network. These dHb clusters are also preserved during odor stimulation and govern olfactory responses. Finally, we show that functional dHb clusters overlap with genetically defined dHb neurons [4], which regulate experience-dependent fear. Thus, we propose that the dHb is composed of functionally, spatially, and genetically distinct microcircuits that regulate different behavioral programs.

By the end of
today's lecture you
will be able to:

- Motivate the use of dimensionality reduction in neuroscience
- Identify several popular methods and their limitations
- Define what eigenvalues and eigenvectors are and determine them using Python
- Describe the steps of PCA and implement it in Python
 - Define and generate a **covariance matrix**

A non-comprehensive map of techniques commonly used to reduce the dimensionality of neural data.



Independent component analysis (ICA) doesn't enforce an orthogonal principle component, and is therefore useful for sources that may not be orthogonal — for example, voices at a cocktail party

Focus for today (and widely used, tested, understood)

Principal components analysis (PCA)
allows us to find axes in high dimensional space that express the information contained in the data with less variables than the original dataset.

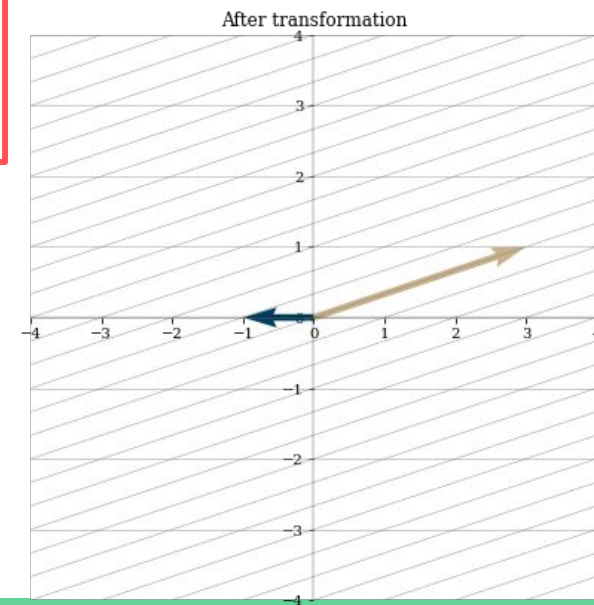
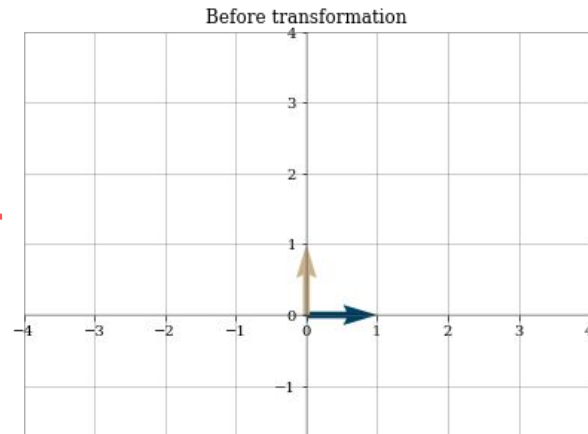
By the end of
today's lecture you
will be able to:

- Motivate the use of dimensionality reduction in neuroscience
- Identify several popular methods and their limitations
- Define what eigenvalues and eigenvectors are and determine them using Python
- Describe the steps of PCA and implement it in Python
 - Define and generate a **covariance matrix**

Reminders

- Matrices apply linear transformations and can be used to change basis vectors
- Every vector has a span, a line that passes through the origin and its tip

$$\begin{bmatrix} -1 & 3 \\ 0 & 1 \end{bmatrix}$$




Eigenvectors

Most vectors will get knocked off their span in a linear transformation
— those that don't are **eigenvectors**

(And so, we can also think about linear transformations in terms of eigenvectors).

Eigenvectors don't
move off their span,
but they can
change length.

Before transformation



After transformation




Image: [NMA](#)

Eigenvectors

Most vectors will get knocked off their span in a linear transformation — those that don't are **eigenvectors**

(And so, we can also think about linear transformations in terms of eigenvectors).

Eigenvectors don't move off their span, but they can change length.

The **eigenvalue** is a **scalar** that tells you how much!

$$W\mathbf{v} = \lambda\mathbf{v}$$


Matrix x Eigenvector = Eigenvalue x Eigenvector

Solve for lambda (λ) and \mathbf{v}

Image: [NMA](#); another explanation:
<https://www.youtube.com/watch?v=PFDu9oVAE-g>

By the end of
today's lecture you
will be able to:

- Motivate the use of dimensionality reduction in neuroscience
- Identify several popular methods and their limitations
- Define what eigenvalues and eigenvectors are and determine them using Python
- Describe the steps of PCA and implement it in Python
 - Define and generate a **covariance matrix**



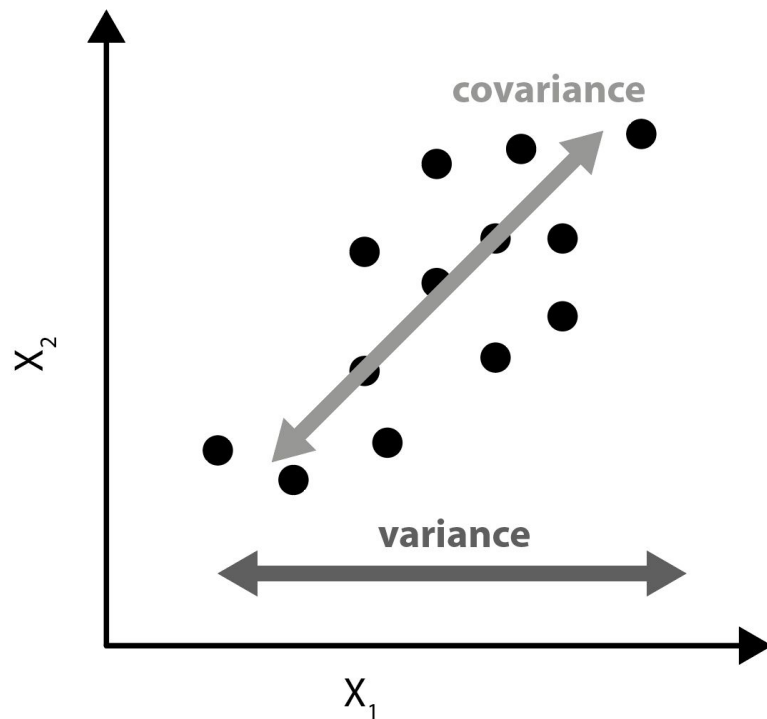
The arrow captures the direction where the **variance is maximized**. Points that are furthest away in the first panel are *still* the furthest away.

We can take two variables (age & height),
and describe them by one vector.

Performing PCA involves five steps


(0) Subtract the mean (and optionally, normalize)

1. Calculate the **covariance matrix** between the variables.



Covariance matrix (Σ)

$\Sigma =$




Variance on diagonal,
covariance on
off-diagonal

Symmetric!

$$\Sigma_{i,j} = \Sigma_{j,i}$$

The **covariance matrix** helps us find the dimension where the variance is maximal

- The **eigenvectors** of the covariance matrix are the Principal Components — the direction which the data is stretched
- The **eigenvalues** quantify the amount of variance explained by each Principal Component



$$x1^2 = x^2 + y^2$$

$$x1 = \sqrt{\text{age}^2 + \text{height}^2}$$

Performing PCA involves five steps

1. Subtract the mean (and optionally, normalize)
2. Calculate the **covariance matrix** between the variables.
3. Extracting the factors by **rotation** (eigenvector decomposition).
4. Projecting mean-centered data onto the eigenvectors, using matrix multiplication



$$\mathbf{S} = \mathbf{X} \mathbf{W}$$

The diagram illustrates the matrix equation $\mathbf{S} = \mathbf{X} \mathbf{W}$ for Principal Component Analysis. It consists of three rectangular boxes: the first box on the left contains the bold letter **S**, the middle box contains the bold letter **X**, and the third box on the right contains the bold letter **W**. An equals sign (=) is placed between the first and second boxes, and between the second and third boxes. Below the first box, an upward-pointing arrow is positioned above the text "scores". Below this, the text "latent variables" is centered. Below the third box, an upward-pointing arrow is positioned above the text "weights". Below this, the text "coefficients" is centered, and at the bottom, the text "loadings" is centered.

↑
"scores"
"latent variables"

↑
"weights"
"coefficients"
"loadings"

How do we decide how many components are enough?



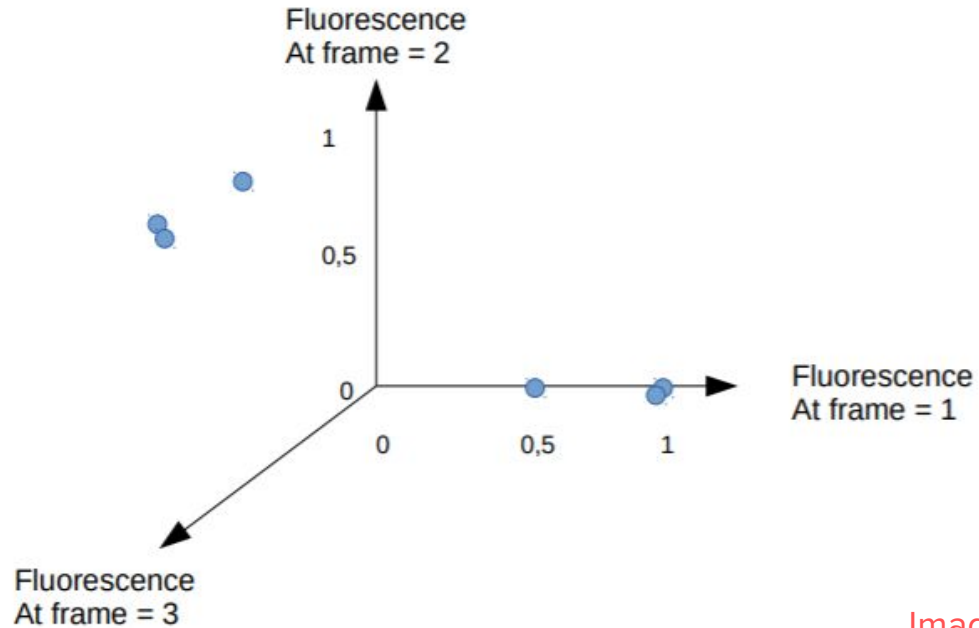
Performing PCA involves five steps

1. Subtract the mean (and optionally, normalize)
2. Calculate the **covariance matrix** between the variables.
3. Extracting the factors by **rotation** (eigenvector decomposition).
4. Projecting mean-centered data onto the eigenvectors
5. Interpreting:
 - a. Determining the **number** of factors needed
 - b. Interpreting the **meaning** of factors.

(we'll work through these steps in the notebook)

Traces =
(6 x 100)

3 time frames		
0	1	0,5
0	1	1
0	1	1
1	0	0
0,5	0	0
1	0	0



[Image source](#)

We can also perform PCA on *time series*

Here, there are 6 neurons recorded over 3 time frames, creating 6 data points in a 3D space. (In reality, there are more like 54000 recorded frames, so this plot should have 50,000 dimensions)

Dataset for today

First author:
**Cynthia
Chestek**

(Read about her
work!)



10742 • The Journal of Neuroscience, October 3, 2007 • 27(40):10742–10750

Behavioral/Systems/Cognitive

Single-Neuron Stability during Repeated Reaching in Macaque Premotor Cortex

Cynthia A. Chestek,^{1*} Aaron P. Batista,^{1,2*} Gopal Santhanam,¹ Byron M. Yu,¹ Afsheen Afshar,^{1,3} John P. Cunningham,¹
Vikash Gilja,⁴ Stephen I. Ryu,^{1,5} Mark M. Churchland,^{1,2} and Krishna V. Shenoy^{1,2}

¹Department of Electrical Engineering, ²Neurosciences Program, ³Medical Scientists Training Program, ⁴Department of Computer Science, and

⁵Department of Neurosurgery, Stanford University, Stanford, California 94305

PI: **Krishna Shenoy**

(Read about his life
& work)



<https://www.jneurosci.org/content/27/40/10742>

Resources

Python Data Science Handbook, [In Depth: PCA](#) & [In-Depth K-Means](#)

StatQuest PCA [overview](#) and in [more detail](#).

Dimensionality Reduction, *Neural Data Science* Chapter 8

[Dimensionality reduction for large-scale neural recordings | Nature Neuroscience](#)

[An introduction to machine learning with scikit-learn](#)

[Steve Brunton's SVD Lecture Series](#)